



COMP 5300 / COMP 4600

Deep Learning for NLP

Lecture

Administrivia

- The class will be taught in two sessions:
 - Mon 2 – 3:15 pm Dandeneau 220
 - Thu 3:30 – 4:45 pm Olsen 405
- Homeworks, slides, topics:
 - class website:
<https://text-machine-lab.github.io/dl4nlp-s2023/>
- Class announcements, questions:
 - Make sure to join the class Discord server!
 - Invite link on Blackboard!

Plan for next week

- Homework #1
 - Will be posted tonight
 - Due next Thu 2/2/23 before class!
 - Use colab to do the homework
 - Colab tutorial: <https://neptune.ai/blog/how-to-use-google-colab-for-deep-learning-complete-tutorial>
 - Mon 1/30/23
 - Lecture, cont'd: Text classification and lexical embeddings
 - Ask questions about HW1
 - Thu 2/2/23
 - Live coding session for word2vec homework (bring your laptops!)
 - Whatever's not finished will become your Homework #2
 - HW1 due; HW2 (w2v) assigned
-

Lecture outline

- Lexical embeddings (= word vectors)
 - count-based (sparse, dense)
 - prediction-based (neural embeddings)
 - evaluating lexical embeddings
- Tokenization
- Text classification



Lexical embeddings

vector space representation for word-level semantics

Naive Representation of Text Documents

One-hot representation for each word:

[0 0 0 0 1 0 0 0 0]

– dimensionality is $|V|$, size of your chosen vocabulary

Compare two documents, e.g. for classification

[0 1 0 0 1 1 0 0 1 0]

[0 0 0 0 1 0 0 0 1 0]

- 1's in positions corresponding to the words present in the document
- could be (normalized) counts instead.

Use e.g. cosine or set-membership similarity measures to compute “distance” between two documents

Count-based Vectors

Can represent a document as a “bag of words”

Turn your text into a vector of word counts

For example, a movie review:

An **unpleasant**, **humorless** slog through the **muck** of **low-budget** January horror **fodder** that is neither frightening or particularly entertaining, **blandly ambling** from **tired** jump-scare to **tired** jump-scare."

`<unpleasant: 1, the: 23, dog: 0, tired:3, ...>`

Bag-of-Words (BOW) Representation

Documents

We will now have small
Gives a function based
Using entropy of traffic
We study the complexity
of influencing elections
through bribery. How
computationally complex
is it for an external actor
to determine whether by
a certain amount of
bribing voters a specified
candidate can be made
the election's winner? We
study this problem for
election systems as varied
as scoring ...



Vector-space representation

	D1	D2	D3	D4	D5
complexity	2		3	2	3
algorithm	3			4	4
entropy	1			2	
traffic		2	3		
network		1	4		

Term-document matrix

Bag-of-Words (BOW) Representation

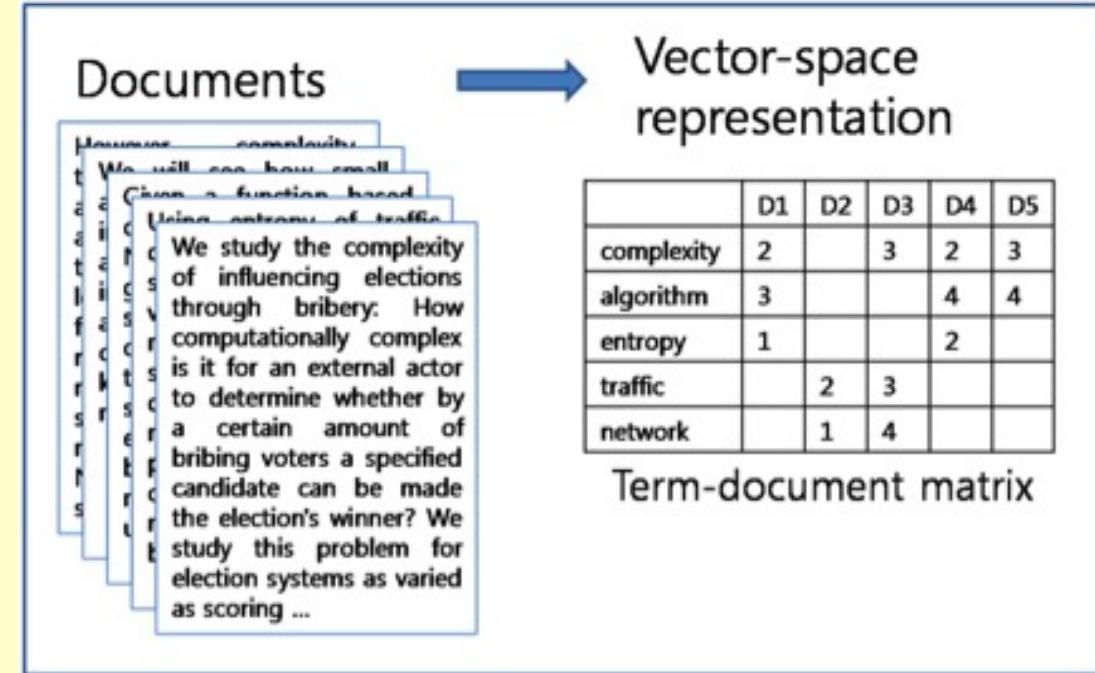
Counts could be normalized by frequency

Probability (divide by corpus size)

Conditional probability (divide by co-occurrence count)

Pointwise mutual information

TF-IDF



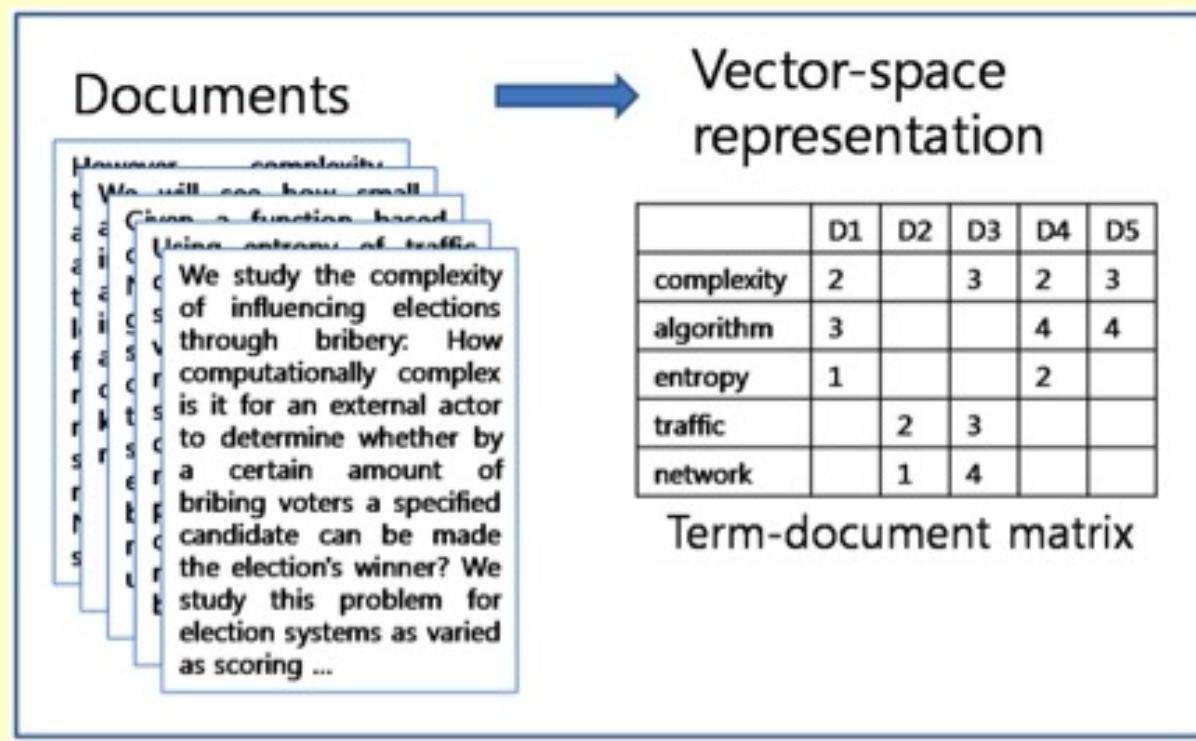
BOW representation works well for tasks invariant to word order, such as text classification

TF-IDF Weighting

$TF-IDF(x) = \text{term frequency in document} / \# \text{ of documents in which the term occurs.}$

IDF ("inverse document frequency") penalizes counts (= reduces weights) of words that occur in many documents

A document could be a sentence, a paragraph, a chapter, a movie/product review, etc.

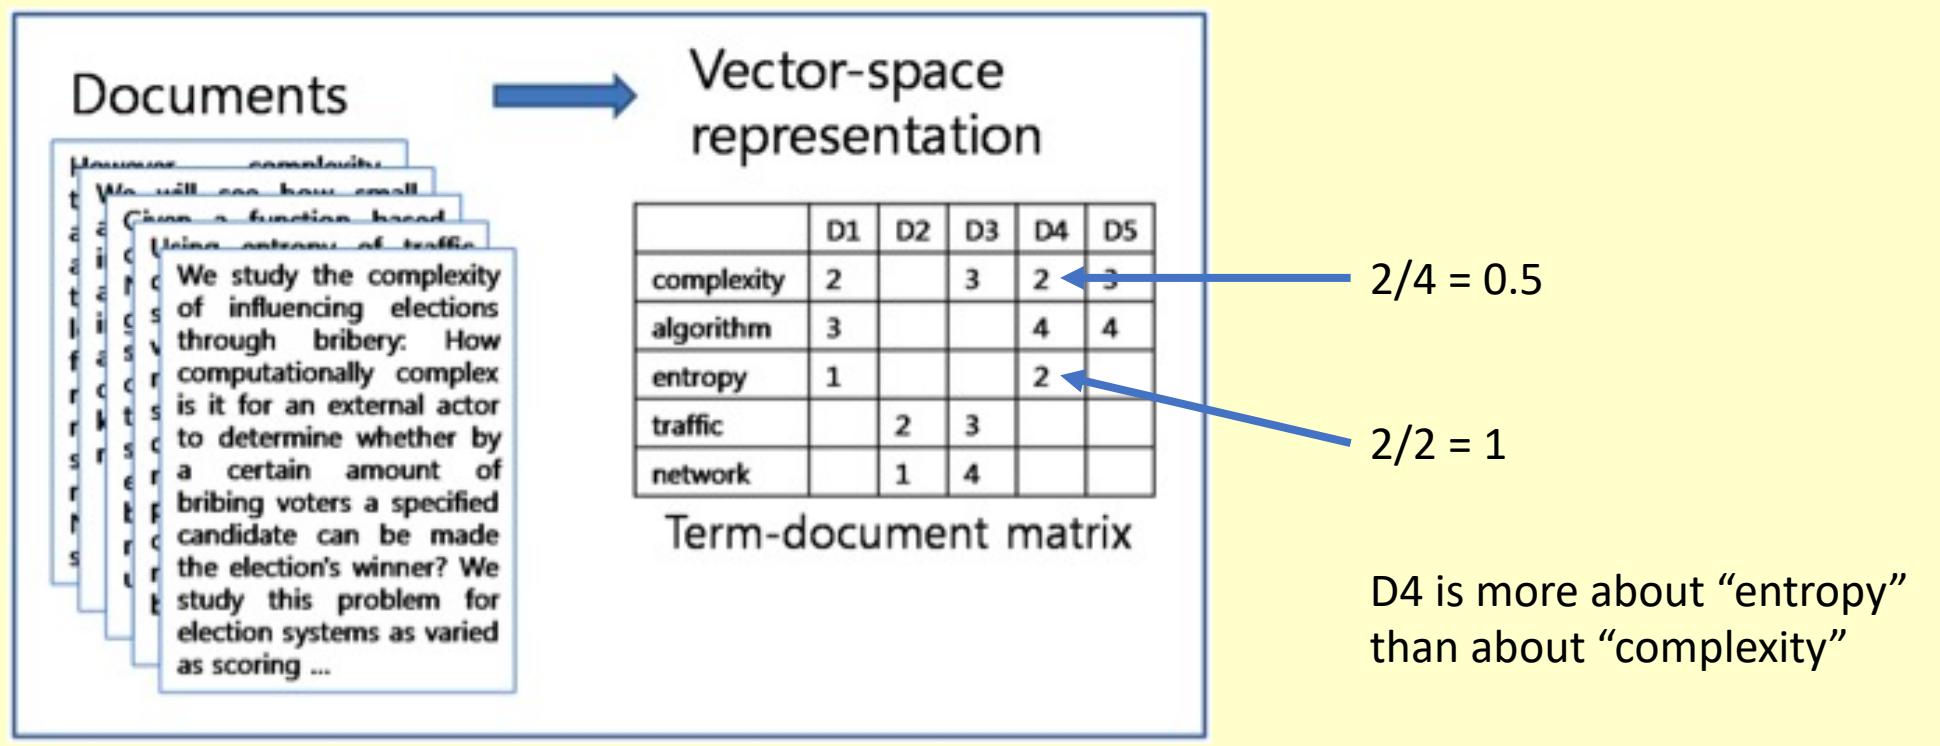


TF-IDF Weighting

$TF-IDF(x) = \text{term frequency in document} / \# \text{ of documents in which the term occurs.}$

IDF (“inverse document frequency”) penalizes counts (= reduces weights) of words that occur in many documents

A document could be a sentence, a paragraph, a chapter, a movie/product review, etc.



Similarity Measures

$$Dice(A, B) = \frac{|A \cap B|}{\frac{1}{2}(|A| + |B|)}; \quad Dice^\dagger(\vec{X}, \vec{Y}) = \frac{\sum_i \min(x_i, y_i)}{\frac{1}{2}(\sum_i x_i + \sum_i y_i)}$$

$$Jaccard(A, B) = \frac{|A \cap B|}{|A \cup B|}; \quad Jaccard^\dagger(\vec{X}, \vec{Y}) = \frac{\sum_i \min(x_i, y_i)}{\sum_i \max(x_i, y_i)}$$

$$\cos(\vec{X}, \vec{Y}) = \frac{\vec{X} \cdot \vec{Y}}{|\vec{X}| |\vec{Y}|} = \frac{\sum_i x_i y_i}{\sqrt{\sum_i x_i^2} \sqrt{\sum_i y_i^2}}$$

$$Euclidean-Distance(\vec{X}, \vec{Y}) = |\vec{X} - \vec{Y}| = \sqrt{\sum_i (x_i - y_i)^2}$$

$$L_1 \text{ norm} = \sum_i |x_i - y_i| = 2(1 - \sum_i \min(x_i, y_i))$$

$$D(p||q) = \sum_i p_i \log \frac{p_i}{q_i}$$

$$JS(p||q) = \frac{1}{2} \left[D(p||\frac{p+q}{2}) + D(q||\frac{p+q}{2}) \right]$$

$$\alpha\text{-}skew(p, q) = D(p||\alpha \cdot q + (1 - \alpha) \cdot p)$$

Zipf's law – representations are always sparse

Zipf's Law (long tail phenomenon):

The frequency of any word is inversely proportional to its rank in the frequency table

A large number of events occur with low frequency. You might have to wait an arbitrarily long time to get valid statistics on low frequency events



Words with frequency
of less than one in 50,000
make up 20-30% of
newswire reports
(Dunning, 1993)

What about words?

One-hot representation for each word:

[0 0 0 0 1 0 0 0 0]

– dimensionality is $|V|$, size of your chosen vocabulary

Can we do better?

Representing words

One-hot representation for each word:

[0 0 0 0 1 0 0 0 0]

– dimensionality is $|V|$, size of your chosen vocabulary

Can we do better?

Can we make it so that similar words have similar representations?

Distributional hypothesis

- Meaning as an invariant of similar words (Mel'chuk)
- Similar words are used in similar contexts.
- This is known as the “distributional hypothesis” (Harris, 1985), or the “strong contextual hypothesis” (Miller and Charles, 1991), and related to the much-quoted remark by Firth (1957)

“You shall know a word by the company it keeps”

Meaning from context

Consider occurrence contexts of an unknown word, **tezgüino**

C1: A bottle of **tezgüino** is on the table

C2: Everybody likes **tezgüino**

C3: Don't have **tezgüino** before you drive.

C4: We make **tezgüino** out of corn.

Distributional statistics for **tezgüino**:

	C1	C2	C3	C4	...
<i>tezgüino</i>	1	1	1	1	
<i>loud</i>	0	0	0	0	
<i>motor oil</i>	1	0	0	1	
<i>tortillas</i>	0	1	0	1	
<i>choices</i>	0	1	0	0	
<i>wine</i>	1	1	1	0	

What about words?

Can we make it so that similar words have similar representations?

- Represent each word as a collection of contexts in which it has occurred in a corpus of texts

	C1	C2	C3	C4	...
<i>tezgüino</i>	1	1	1	1	
<i>loud</i>	0	0	0	0	
<i>motor oil</i>	1	0	0	1	
<i>tortillas</i>	0	1	0	1	
<i>choices</i>	0	1	0	0	
<i>wine</i>	1	1	1	0	

Concordance for “showed”

KWIC concordance (Key Word In Context)

1 could find a target. The librarian
2 elights in. The young lady teachers
3 ingly. The young gentlemen teachers
4 seeming vexation). The little girls
5 n various ways, and the little boys
6 t genuwyne?" Tom lifted his lip and
7 is little finger for a pen. Then he
8 ow's face was haggard, and his eyes
9 not overlook the fact that Tom even
10 own. Two or three glimmering lights
11 ird flash turned night into day and
12 that grew about their feet. And it
13 he first thing his aunt said to him
14 p from her lethargy of distress and
15 ent a new burst of grief from Becky
16 shudder quiver all through him. He

“showed off” - running hither and thither w
“showed off” - bending sweetly over pupils
“showed off” with small scoldings and other
“showed off” in various ways, and the littl
“showed off” with such diligence that the a
showed the vacancy. “Well, all right,” sai
Huckleberry how to make an H and an
the fear that was upon him. When he
a marked aversion to these inquests
where it lay, peacefully sleeping,
every little grass-blade, separate
three white, startled faces, too. A
him that he had brought his sorrows
good interest in the proceedings. S
Tom that the thing in his mind had
Huck the fragment of candle-wick pe

Concordance for “showed”

KWIC concordance (Key Word In Context)

1 could find a target. The librarian
2 elights in. The young lady teachers
3 ingly. The young gentlemen teachers
4 seeming vexation). The little girls
5 n various ways, and the little boys
6 t genuwyne?" Tom lifted his lip and
7 is little finger for a pen. Then he
8 ow's face was haggard, and his eyes
9 not overlook the fact that Tom even
10 own. Two or three glimmering lights
11 ird flash turned night into day and
12 that grew about their feet. And it
13 he first thing his aunt said to him
14 p from her lethargy of distress and
15 ent a new burst of grief from Becky
16 shudder quiver all through him. He

“showed off” - running hither and thither w
“showed off” - bending sweetly over pupils
“showed off” with small scoldings and other
“showed off” in various ways, and the littl
“showed off” with such diligence that the a
showed the vacancy. “Well, all right,” sai
showed Huckleberry how to make an H and an
showed the fear that was upon him. When he
showed a marked aversion to these inquests
showed where it lay, peacefully sleeping,
showed every little grass-blade, separate
showed three white, startled faces, too. A
showed him that he had brought his sorrows
showed good interest in the proceedings. S
showed Tom that the thing in his mind had
showed Huck the fragment of candle-wick pe

Concordance for “showed”

KWIC concordance (Key Word In Context)

1 could find a target. The librarian
2 elights in. The young lady teachers
3 ingly. The young gentlemen teachers
4 seeming vexation). The little girls
5 n various ways, and the little boys
6 t genuwyne?" Tom lifted his lip and
7 is little finger for a pen. Then he
8 ow's face was haggard, and his eyes
9 not overlook the fact that Tom even
10 own. Two or three glimmering lights
11 ird flash turned night into day and
12 that grew about their feet. And it
13 he first thing his aunt said to him
14 p from her lethargy of distress and
15 ent a new burst of grief from Becky
16 shudder quiver all through him. He

“showed off” - running hither and thither w
“showed off” - bending sweetly over pupils
“showed off” with small scoldings and other
“showed off” in various ways, and the littl
“showed off” with such diligence that the a
showed the vacancy. “Well, all right,” sai
showed Huckleberry how to make an H and an
showed the fear that was upon him. When he
showed a marked aversion to these inquests
showed where it lay, peacefully sleeping,
showed every little grass-blade, separate
showed three white, startled faces, too. A
showed him that he had brought his sorrows
showed good interest in the proceedings. S
showed Tom that the thing in his mind had
showed Huck the fragment of candle-wick pe

Context definition

Example: “was haggard and his eyes **showed** the fear that was upon”

Window BOW context (+5/-5):

{was, haggard, and, his, eyes, the, fear, that, was, upon}

Structured context (+5/-5):

$\{(was, -5), (haggard, -4), (and, -3), (his, -2), (eyes, -1), (the, +1), (fear, +2), (that, +3), (was, +4), (upon, +5)\}$

Dependency context:

$\{(eyes, \text{NSUBJ}), (\text{fear}, \text{DOBJ}), \dots\}$

develop bnc freq = 23388

<u>object</u>	<u>14323</u> 5.5	<u>subject</u>	<u>4512</u> 3.5	<u>modifier</u>	<u>3153</u> 2.4	<u>and/or</u>	<u>960</u> 0.5
skill	432 38.26	compact	10 19.46	further	192 46.65	grow	45 27.52
technique	219 30.08	patient	58 18.43	rapidly	97 40.95	implement	26 26.95
strategy	164 27.61	situation	44 16.53	jointly	51 39.89	Maintain	39 25.87
model	195 24.98	capitalism	14 16.51	fully	116 38.68	refine	12 24.94
idea	223 24.14	embryo	11 16.43	originally	71 37.65	expand	21 24.78
method	168 23.77	scientist	22 16.11	well	132 31.64	market	15 24.29
theory	158 23.45	Freud	10 15.47	independently	28 30.91	test	19 21.03
system	439 23.39	relationship	42 15.2	subsequently	39 30.3	sustain	14 20.7
understanding	103 23.24	Scotvec	8 15.17	specially	31 29.26	research	10 20.11
technology	139 22.69	Ici	9 14.72	later	59 29.21	explore	15 19.84
expertise	50 21.69	Aea	6 14.33	also	184 28.96	extend	20 19.28
methodology	29 21.33	Ibm	14 14.19	highly	59 28.58	improve	17 17.74
concept	89 21.17	company	61 12.94	quickly	51 27.9	Promote	14 17.42
relationship	144 20.67	friendship	10 12.38	gradually	31 27.33	train	12 16.86

Bag-of-Words (BOW) Representation

Counts could be normalized:

Conditional probability (divide by co-occurrence count)

TF-IDF

Pointwise mutual information

Count (x) = how many times a term x occurs within a context window of the target word w ?

Conditional_Probability(x) = divide *Count*(x) by how often many times the target word w occurs in our corpus of texts.

TF-IDF (x) = term frequency in document / # of documents in which the term occurs.

$$\text{pmi}(x; y) \equiv \log \frac{p(x, y)}{p(x)p(y)} = \log \frac{p(x|y)}{p(x)} = \log \frac{p(y|x)}{p(y)}$$



develop bnc freq = 23388

<u>object</u>	<u>14323</u> 5.5	<u>subject</u>	<u>4512</u> 3.5	<u>modifier</u>	<u>3153</u> 2.4	<u>and/or</u>	<u>960</u> 0.5
skill	432 38.26	compact	10 19.46	further	192 46.65	grow	45 27.52
technique	219 30.08	patient	58 18.43	rapidly	97 40.95	implement	26 26.95
strategy	164 27.61	situation	44 16.53	jointly	51 39.89	Maintain	39 25.87
model	195 24.98	capitalism	14 16.51	fully	116 38.68	refine	12 24.94
idea	223 24.14	embryo	11 16.43	originally	71 37.65	expand	21 24.78
method	168 23.77	scientist	22 16.11	well	132 31.64	market	15 24.29
theory	158 23.45	Freud	10 15.47	independently	28 30.91	test	19 21.03
system	439 23.39	relationship	42 15.2	subsequently	39 30.3	sustain	14 20.7
understanding	103 23.24	Scotvec	8 15.17	specially	31 29.26	research	10 20.11
technology	139 22.69	Ici	9 14.72	later	59 29.21	explore	15 19.84
expertise	50 21.69	Aea	6 14.33	also	184 28.96	extend	20 19.28
methodology	29 21.33	Ibm	14 14.19	highly	59 28.58	improve	17 17.74
concept	89 21.17	company	61 12.94	quickly	51 27.9	Promote	14 17.42
relationship	144 20.67	friendship	10 12.38	gradually	31 27.33	train	12 16.86

Word meaning represented by word embeddings



Count-based frequency vectors

Dimensionality of vocabulary $|V|$, 100K

Sparse

Prediction-based continuous “dense” vectors

Lower dimensionality (commonly 200-500)

Continuous





Dense word embeddings

Dense Word Embeddings

1. Reduce dimensionality of count-based representation
 - Principal Component Analysis (PCA), singular value decomposition (SVD) (also “Latent Semantic Analysis”, LSA).
2. Learn embeddings as parameters in a learning task, where a cost function is tied to the context
 - Different approximations to the log likelihood of the full corpus

Singular Value Decomposition

Mathematical apparatus:

- ◆ $\text{rank}(A)$
 - number of linearly independent columns in $A_{m \times n}: \mathbb{R}^n \rightarrow \mathbb{R}^m$
 - linear transform of rank r maps the basis vectors of the pre-image into r linearly independent basis vectors of the image
- ◆ Singular Value Decomposition
 - Matrix A can be represented as
 - » $A = U \Sigma V^T$ where
 - columns of U and V are left and right eigenvectors of $A A^T$
 - U and V orthogonal ($VV^T=I$) , and Σ is a diagonal matrix:
 $\Sigma = \text{diag}(\sigma_1, \dots, \sigma_r)$, where σ_i are nonnegative square roots of the r eigenvalues of $A A^T$,

SVD produces a k-rank approximation \hat{A} to matrix A minimizing the “distance” between the two matrices in the form of Frobenius norm (aka 2-norm, Euclidean norm) is minimized:

- ◆ $U_{m \times k}, \Sigma_{k \times k}, V_{k \times n}^T$
- ◆ $\|A - A_k\|_F^2$

$$\|A\|_F \equiv \sqrt{\sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^2}$$

Minimize the objective:

$$\text{Min } \|A - U \Sigma V^T\|_F$$



As linear regression can be interpreted as collapsing a two-dimensional space onto a one-dimensional line, SVD can be thought of as projecting an n-dimensional space onto a k-dimensional space where $n \gg k$.

Folding new count-based vectors into the reduced space:

$$A = U \Sigma V^T$$

$$\rightarrow U^T A = U^T U \Sigma V^T$$

$$\rightarrow U^T A = \Sigma V^T$$

U, V are orthonormal (column vectors are unit length and orthogonal, so $U^T U = I$)

Project a new count-based vector into k-dimensional space

$$a_k = U^T a_{new}$$

$U = (\text{W})\text{o}rds$, $V = (\text{C})\text{o}n\text{tex}ts$

A

U

Σ

V

$$\begin{bmatrix} X \\ |V| \times |V| \end{bmatrix} = \begin{bmatrix} W \\ |V| \times k \end{bmatrix} \begin{bmatrix} \sigma_1 & 0 & 0 & \dots & 0 \\ 0 & \sigma_2 & 0 & \dots & 0 \\ 0 & 0 & \sigma_3 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & \sigma_k \end{bmatrix} \begin{bmatrix} C \\ k \times |V| \end{bmatrix}$$

embedding
for
word i

$$\begin{bmatrix} \quad \\ \quad \\ W \\ \quad \end{bmatrix}_{|V| \times k}$$

$$a_k = W^T a_{new}$$



Neural word embeddings

Embedding Models

Learning count based vectors produces an **embedding matrix** in $\mathbb{R}^{|\text{vocab}| \times |\text{context}|}$:

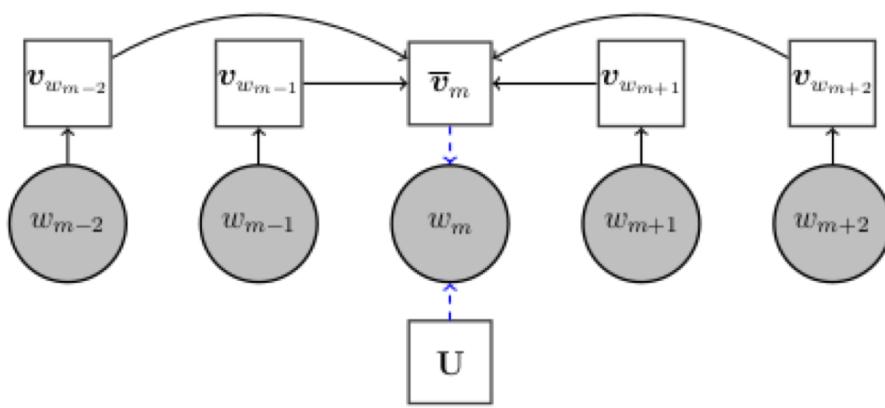
$$\mathbf{E} = \begin{matrix} & \text{bit} & \text{cute} & \text{furry} & \dots \\ \text{kitten} & 0 & 1 & 0 & \dots \\ \text{cat} & 0 & 1 & 1 & \dots \\ \text{dog} & 1 & 0 & 1 & \dots \\ \vdots & \vdots & \vdots & \vdots & \ddots \end{matrix}$$

Rows are word vectors, so we can retrieve them with **one hot vectors** in $\{0,1\}^{|\text{vocab}|}$:

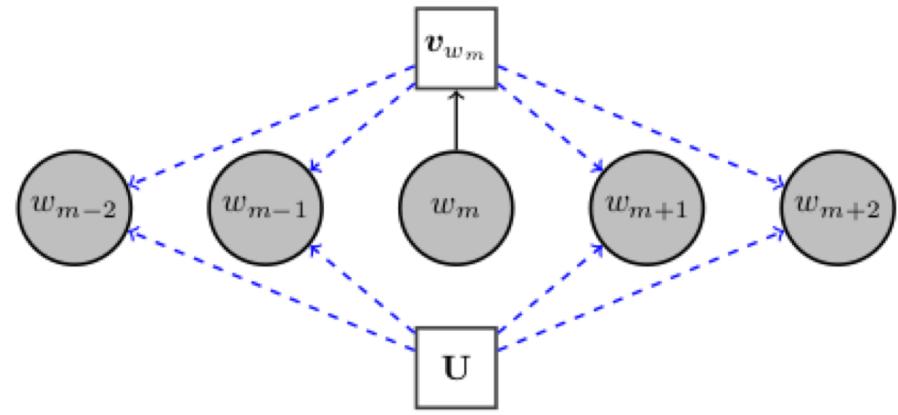
$$\text{onehot}_{\text{cat}} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} \quad \mathbf{cat} = \text{onehot}_{\text{cat}}^\top \mathbf{E}$$

Symbols = unique vectors. Representation = embedding symbols with \mathbf{E} .

word2vec (Mikolov et al 2013)



(a) Continuous bag-of-words (CBOW)



(b) Skipgram

Figure 13.3: The CBOW and skipgram variants of WORD2VEC. The parameter \mathbf{U} is the matrix of word embeddings, and each v_m is the context embedding for word w_m .

word2vec (Mikolov et al 2013)

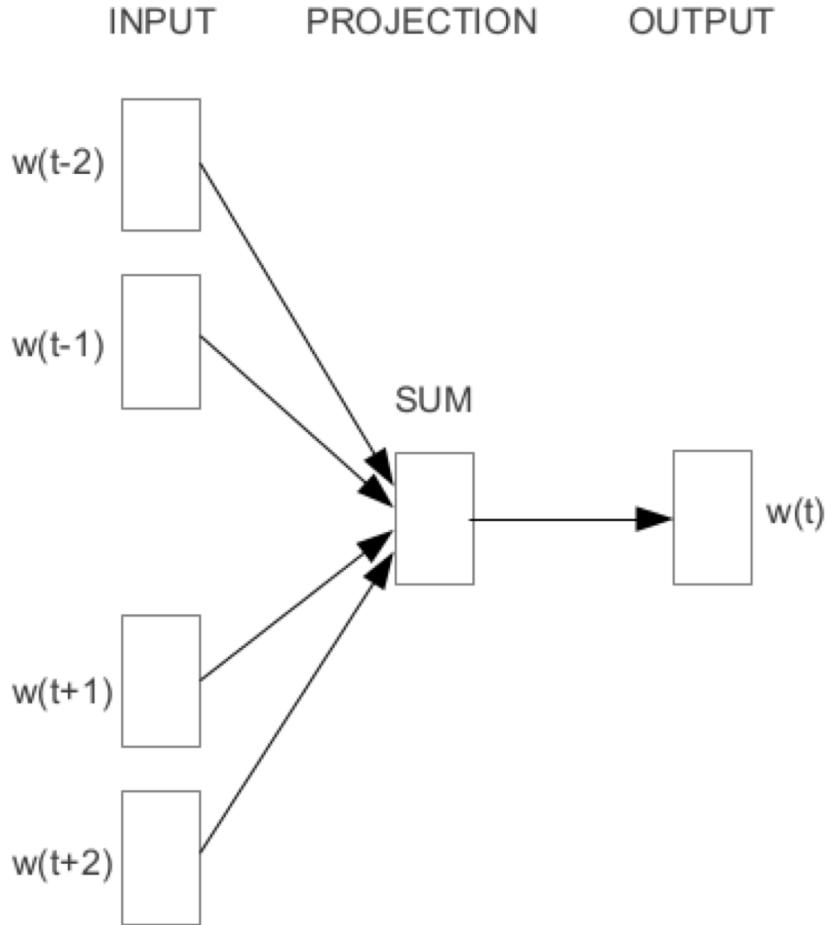
Let vector u_i = the k-dimensional embedding for word i

Let v_j = k-dimensional embedding for context j.

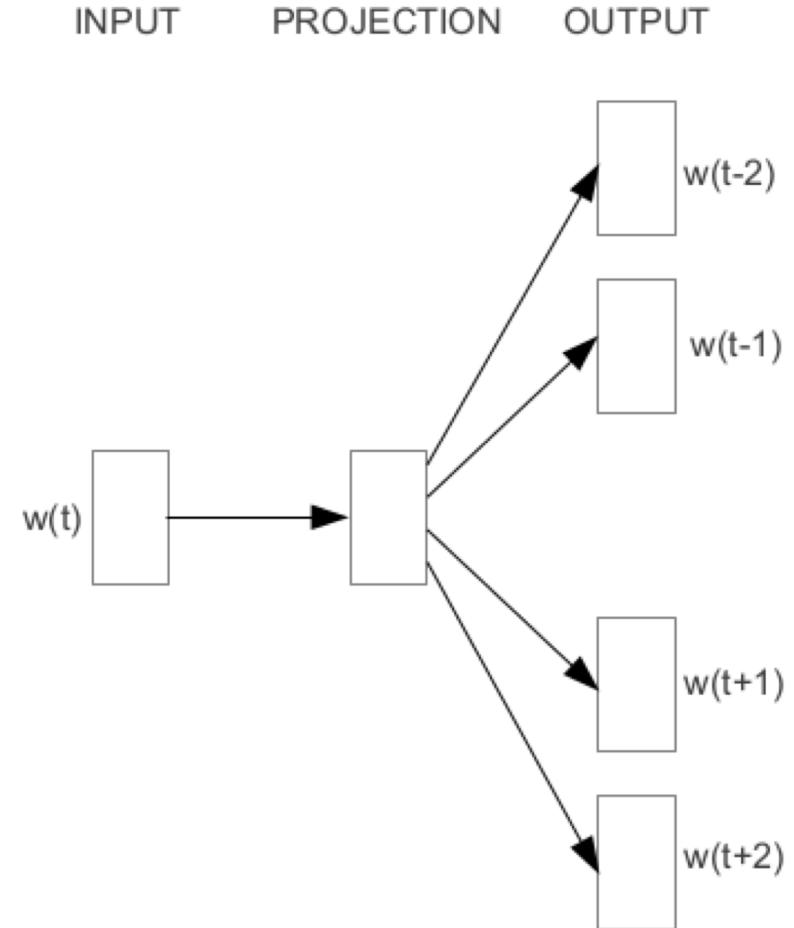
The inner product $u_i \cdot v_j$ represents the compatibility between word i and context j.

By incorporating this inner product into an approximation to the log-likelihood of a corpus, it is possible to estimate both parameters by backpropagation.

word2vec includes two such approximations: continuous bag-of-words (CBOW) and skip-gram.



CBOW



Skip-gram

Word2vec CBOW

BOW b/c order of words doesn't matter

h determines window size

Local context is computed as an average of embeddings
for words in the immediate neighborhood $m - h, m - h + 1, \dots, m + h - 1, m + h$

$$\tilde{v}_m = \frac{1}{2h} \sum_{n=1}^h (v_{m+n} + v_{m-n})$$

Word2vec CBOW

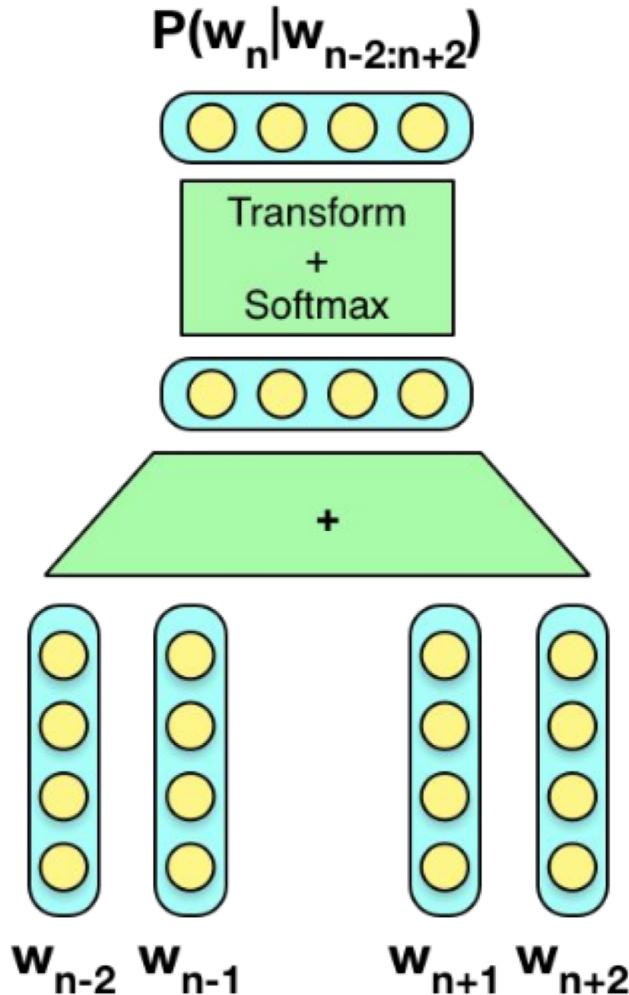
Words are predicted from context

Optimizes approximation to corpus log likelihood

$$\begin{aligned}\log p(\mathbf{w}) &\approx \sum_{m=1}^M \log p(w_m \mid w_{m-h}, w_{m-h+1}, \dots, w_{m+h-1}, w_{m+h}) \\ &= \sum_{m=1}^M \log \frac{\exp(\mathbf{u}_{w_m} \cdot \bar{\mathbf{v}}_m)}{\sum_{j=1}^V \exp_j(\mathbf{u}_i \cdot \bar{\mathbf{v}}_m)} \\ &= \sum_{m=1}^M \mathbf{u}_{w_m} \cdot \bar{\mathbf{v}}_m - \log \sum_{j=1}^V \exp_j(\mathbf{u}_i \cdot \bar{\mathbf{v}}_m).\end{aligned}$$

M is the size of the corpus

word2vec – continuous bag of words (CBOW)



Embed context words. Add them.
Project back to vocabulary size. Softmax.

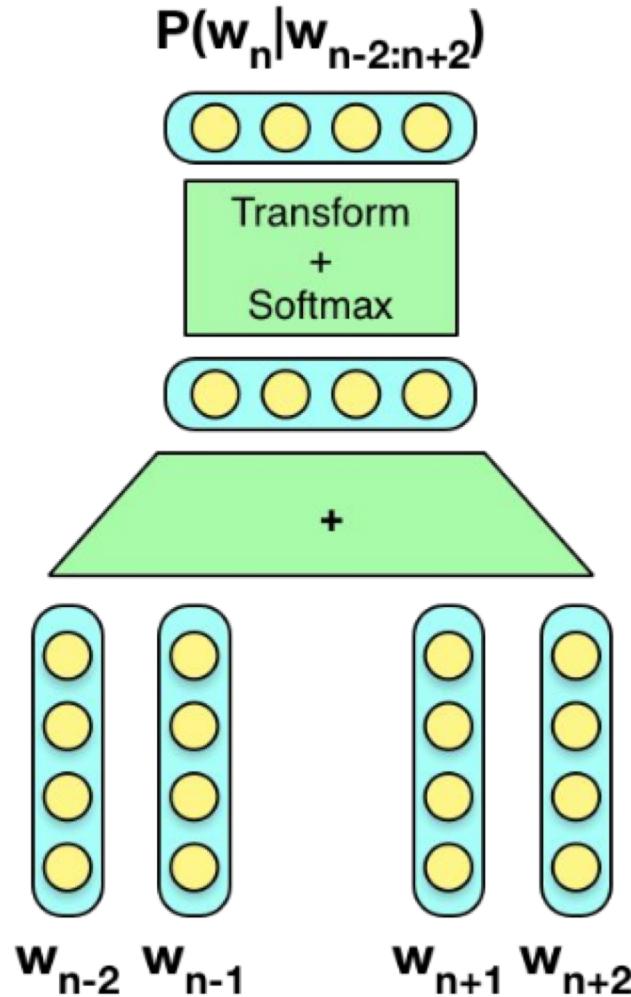
$$\text{softmax}(\mathbf{l})_i = \frac{e^{l_i}}{\sum_j e^{l_j}}$$

$$\begin{aligned} P(t_i | \text{context}(t_i)) &= \text{softmax} \left(\sum_{t_j \in \text{context}(t_i)} \text{onehot}_{t_j}^\top \mathbf{E} W_v \right) \\ &= \text{softmax} \left(\left(\sum_{t_j \in \text{context}(t_i)} \text{onehot}_{t_j}^\top \mathbf{E} \right) W_v \right) \end{aligned}$$

Minimize Negative Log Likelihood:

$$L_{\text{data}} = - \sum_{t_i \in \text{data}} \log P(t_i | \text{context}(t_i))$$

word2vec – continuous bag of words (CBOW)



All linear, so very fast. Basically a cheap way of applying one matrix to all inputs.

Historically, negative sampling used instead of expensive softmax.

NLL minimisation is more stable and is fast enough today.

Variants: position specific matrix per input (Ling et al. 2015).

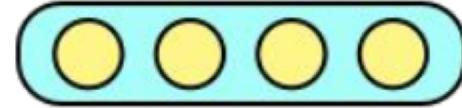
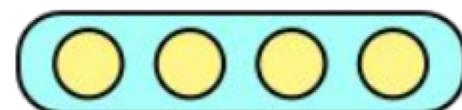
Word2vec Skip-Gram

Contexts are predicted from words

$$\begin{aligned}\log p(\mathbf{w}) &\approx \sum_{m=1}^M \sum_{n=1}^{h_m} \log p(w_{m-n} \mid w_m) + \log p(w_{m+n} \mid w_m) \\&= \sum_{m=1}^M \sum_{n=1}^{h_m} \log \frac{\exp(\mathbf{u}_{w_{m-n}} \cdot \mathbf{v}_{w_m})}{\sum_{j=1}^V \exp(\mathbf{u}_i \cdot \mathbf{v}_{w_m})} + \log \frac{\exp(\mathbf{u}_{w_{m+n}} \cdot \mathbf{v}_{w_m})}{\sum_{j=1}^V \exp(\mathbf{u}_i \cdot \mathbf{v}_{w_m})} \\&= \sum_{m=1}^M \sum_{n=1}^{h_m} \mathbf{u}_{w_{m-n}} \cdot \mathbf{v}_{w_m} + \mathbf{u}_{w_{m+n}} \cdot \mathbf{v}_{w_m} - 2 \log \sum_{j=1}^V \exp(\mathbf{u}_i \cdot \mathbf{v}_{w_m}).\end{aligned}$$

word2vec – Skip-gram

$$\prod_{i=[n-2,n+2] - \{n\}} P(w_i|w_n)$$



w_n

Target word predicts context words.

Embed target word.

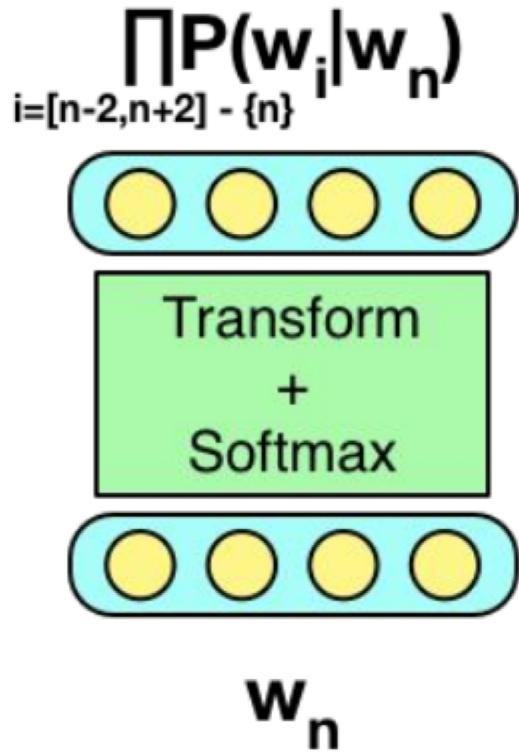
Project into vocabulary. Softmax.

$$P(t_j|t_i) = \text{softmax}(\text{onehot}_{t_i}^\top \mathbf{E} W_v)$$

Learn to estimate likelihood of context words.

$$\begin{aligned} -\log P(\text{context}(t_i)|t_i) &= -\log \prod_{t_j \in \text{context}(t_i)} P(t_j|t_i) \\ &\quad - \sum_{t_j \in \text{context}(t_i)} \log P(t_j|t_i) \end{aligned}$$

word2vec – Skip-gram



Fast: One embedding versus $|C|$ embeddings.

Just read off probabilities from softmax.

- probabilities of the context words, that is

Similar variants to CBoW possible: position specific projections.

Trade off between efficiency and more structured notion of context.

Time Complexity

CBOW and skipgram have a linear time complexity in the size of the word and context representations.

But! they compute a normalized probability over word tokens – a naïve implementation requires summing over the entire vocabulary.

The time complexity of this sum is $O(V \times K)$ for k -dimensional embeddings – which dominates all other computational costs.

One solution is negative **negative sampling**:

Negative sampling is an approximation that eliminates the dependence on vocabulary size!

Negative sampling

Likelihood-based methods are expensive b/c each probability must be normalized over the vocabulary

These probabilities are based on similarity scores between words and contexts for each word in each context.

Can we define an alternative objective based on the same word-context co-occurrence scores $\psi(w, c)$?

Negative sampling

Seek word embeddings that *maximize the difference* between the score for the word observed in context, and the scores for several randomly selected “negative samples” (words that did not occur in that context):

$$\psi(i, j) = \log \sigma(\mathbf{u}_i \cdot \mathbf{v}_j) + \sum_{i' \in \mathcal{W}_{\text{neg}}} \log(1 - \sigma(\mathbf{u}_{i'} \cdot \mathbf{v}_j))$$

- where $\psi(i, j)$ is the score for word i in context j , \mathcal{W}_{neg} is the set of negative samples.

$1 - \sigma(x) = \sigma(-x)$ i.e. probability that x did not occur (cf. $\sigma(x)$ graph)

The objective is to maximize the sum over the corpus:

$$\sum \psi(w_m, c_m)$$

i.e. log product probability of each encountered context in corpus by probability that negative samples for that context didn't occur

Negative Sampling

The set of negative samples W_{neg} is obtained by sampling from a unigram language model.

Unigram language model constructed by exponentiating the empirical word probabilities, setting $\hat{p}(i) \propto (\text{count}(i))^{(3/4)}$.

This has the effect of redistributing probability mass from common to rare words.

The number of negative samples increases the time complexity of training by a constant factor.

5-20 negative samples works for small training sets, and that two to five samples suffice for larger corpora.

Word Embeddings as Matrix factorization

The negative sampling objective is linked to the matrix factorization objective employed in latent semantic analysis.

For a matrix of word-context pairs in which all counts are non-zero, negative sampling is equivalent to factorization of the matrix M,

where $M_{ij} = \text{PMI}(i, j) - \log k$

- each cell in the matrix is equal to the pointwise mutual information of the word and context, shifted by $\log k$, with k equal to the number of negative samples (Levy and Goldberg, 2014)
- k is the number of negative samples in SGNS

Word embeddings are obtained by factoring this matrix with truncated singular value decomposition.

GloVe Embeddings (Stanford)

Another matrix factorization approach. The matrix to be factored is constructed from log co-occurrence counts, $M_{ij} = \text{count}(i, j)$.

Weighted least squares is the objective:

$$\begin{aligned} \min_{\mathbf{u}, \mathbf{v}, b_i, \tilde{b}_j} \quad & \sum_{j=1}^V \sum_{i \in \mathcal{C}} f(M_{ij}) \left(\widehat{\log M_{ij}} - \log M_{ij} \right)^2 \\ \text{s.t.} \quad & \widehat{\log M_{ij}} = \mathbf{u}_i \cdot \mathbf{v}_j + b_i + \tilde{b}_j, \end{aligned}$$

where b_i and \tilde{b}_j are biases for word i and context j , which are estimated jointly with the word embedding \mathbf{u} and context embedding \mathbf{v} .

The weighting function $f(M_{ij})$ is set to 0 at $M_{ij} = 0$

To avoid overweighing frequent context-word pairs:

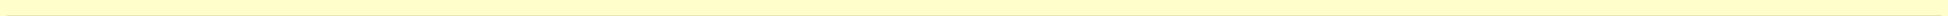
$$f(M_{ij}) = (M_{ij} / \text{threshold})^{3/4} \text{ if } M_{ij} < \text{threshold}, 1 \text{ otherwise}$$



Evaluating lexical embeddings

(vector space representation for word-level semantics)

How good are these representations?



Evaluation of word embeddings

- WordSim-353 (Finkelstein et al. 2003)
contains two sets of English word pairs along with human-assigned similarity judgements
- SimLex-999 (Hill et al. 2016, but has been around since 2014)

Pair	Simlex-999 rating	WordSim-353 rating
<i>coast - shore</i>	9.00	9.10
<i>clothes - closet</i>	1.96	8.00

- Word analogy task (Mikolov et al. 2013), queen = king - man + woman.
- Embedding visualization (nearest neighbors, T-SNE projection)

T-SNE – dimensionality reduction technique

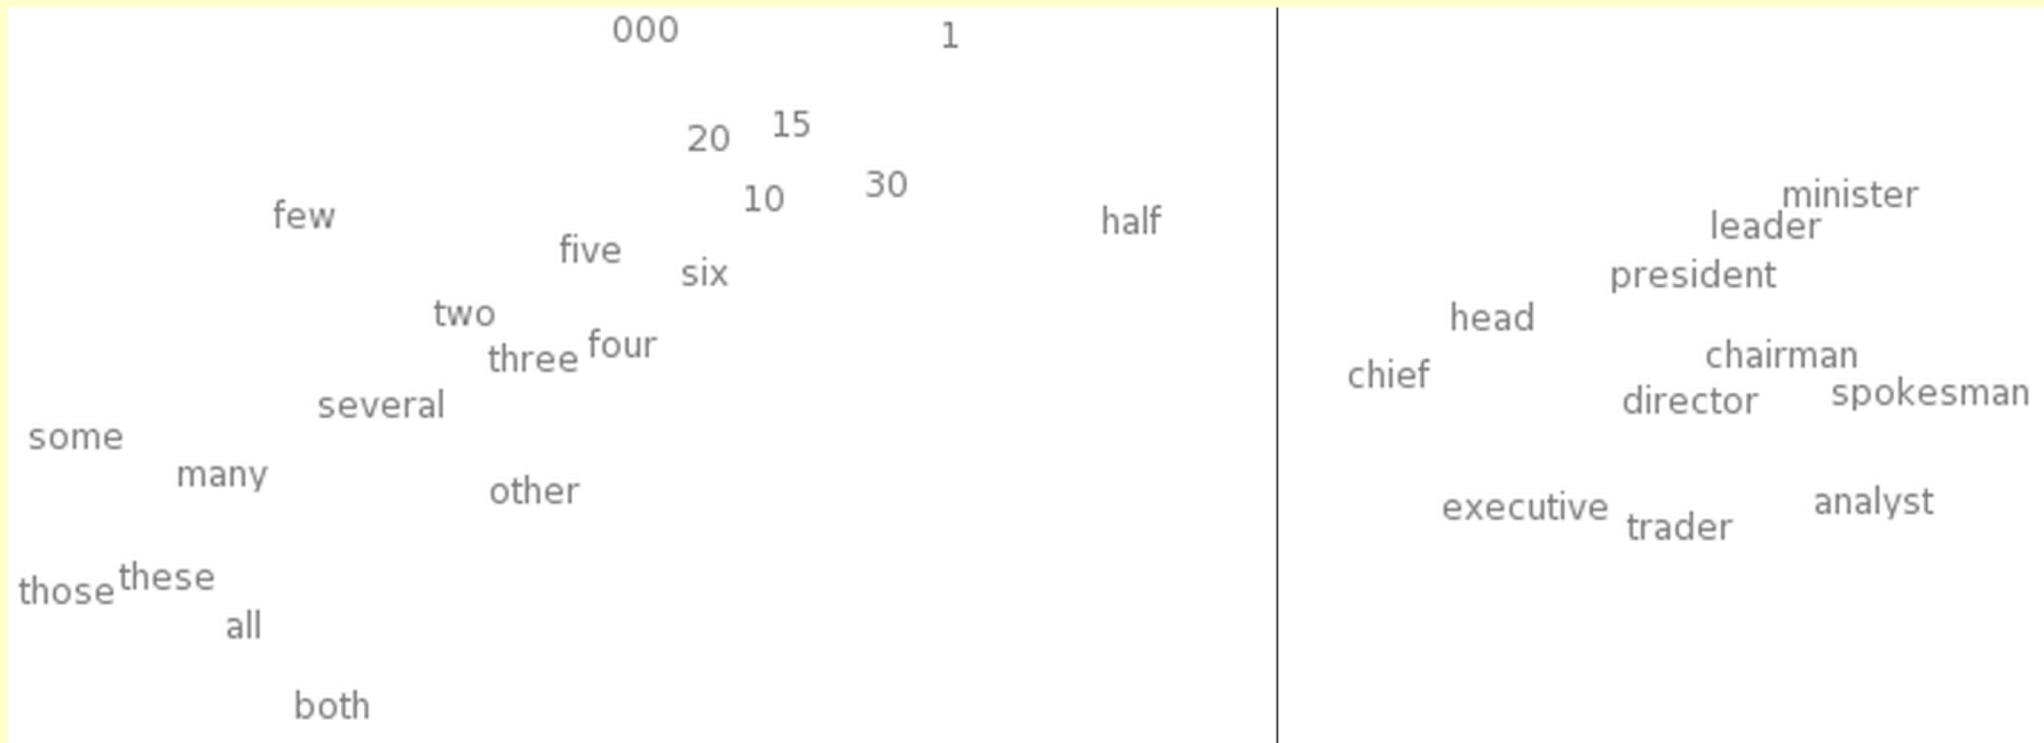
“Distributed stochastic neighbor embedding” (Maaten & Hinton 2008)

Projects points into 2d space by minimizing KL-divergence between two probability distributions between pairs of objects in high-dimensional space, and pairs of objects in low-dimensional space

similar pairs (e.g. via Euclidian distance) have high probability, dissimilar pairs have low probability

Evaluation of word embeddings

Nearest neighbors using t-SNE visualization technique



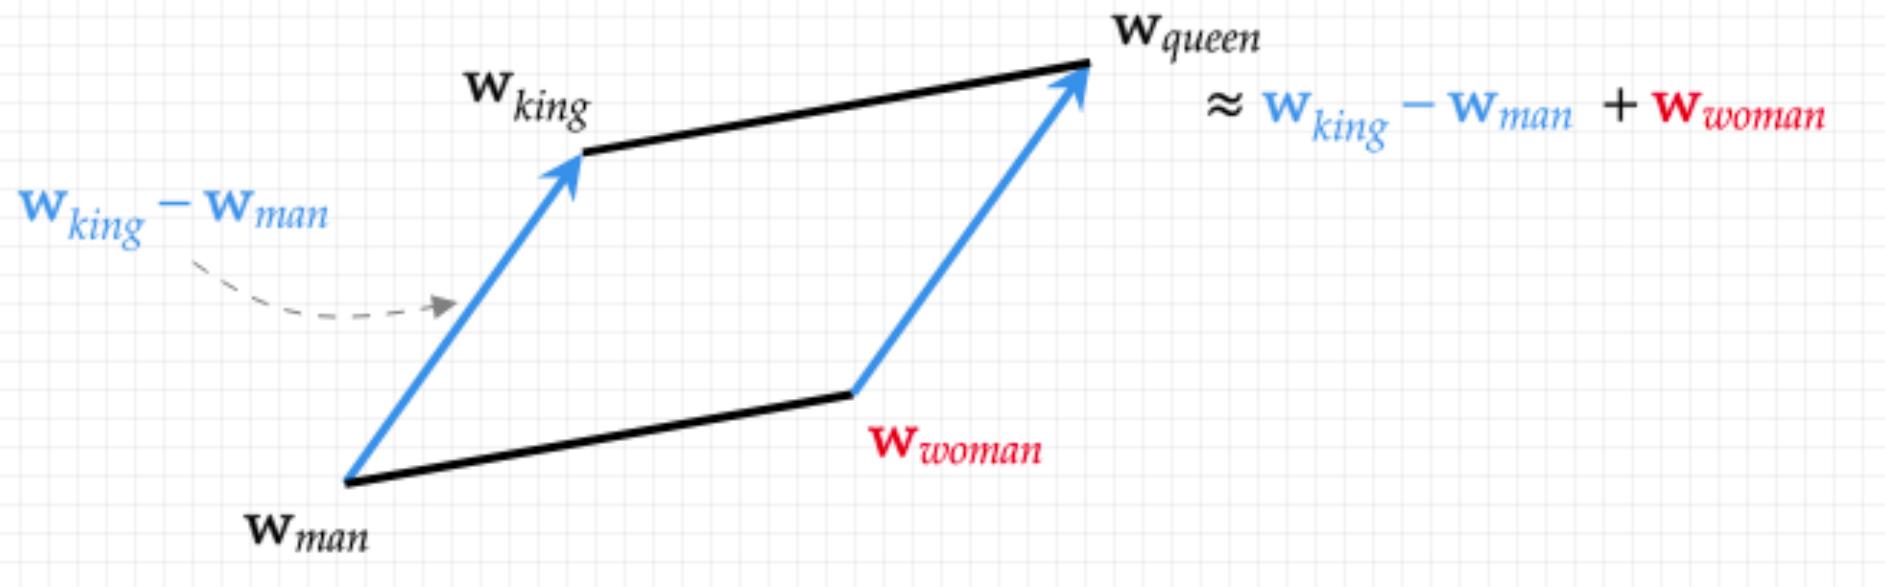
Word Similarity / Relatedness

	Target word	GloVe	SVD
1	phone	telephone	mobile
2	coffee	tea	drinks
3	grammar	vocabulary	grammatical
4	cohesiveness	cohesion	inclusiveness

Table 1: Examples of nearest neighbors in GloVe and SVD

Evaluation of word embeddings

Analogies



Different Ways to do Analogies with the Same Embedding Scheme

ANALOGY

a to b is as a' to b'

TASK: Given a, b, and a', find b'

METHOD 1 (Mikolov et al 2013)

$$b' = \operatorname{argmax}_v (\cos b', b - a + a')$$

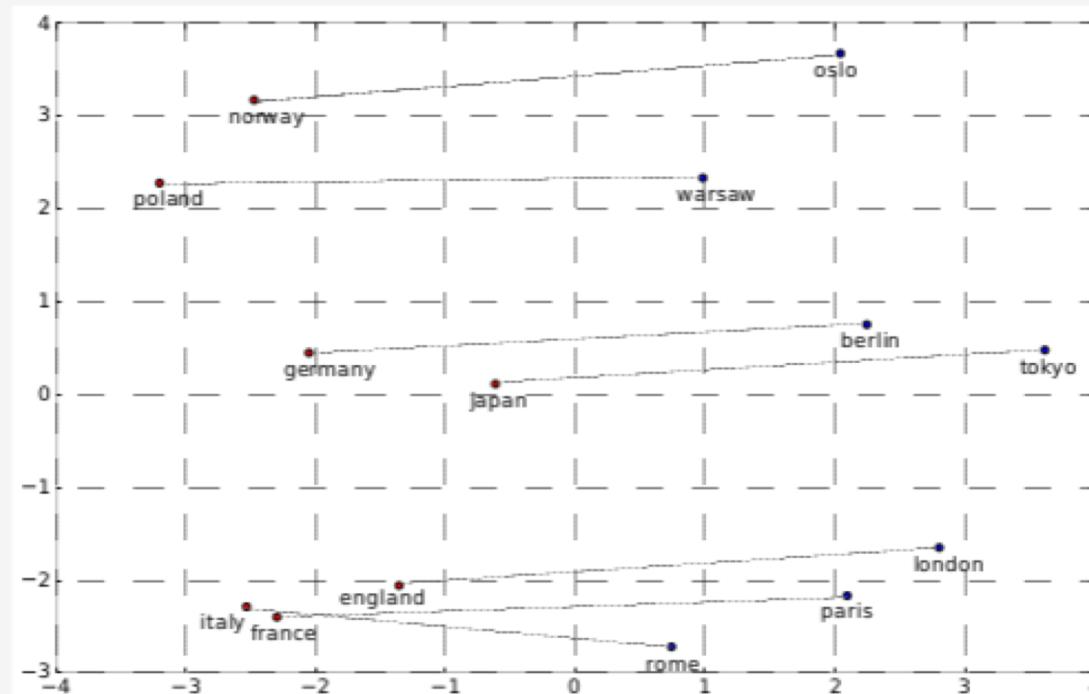
METHOD 2 (Levy & Goldberg 2014)

$$b' = \operatorname{argmax}_v (\cos b' - b, a' - a)$$

The Analogy Test

- Task: given one word, identify its pair with a given linguistic relation.
- Method: solving word analogies: $\overrightarrow{\text{Berlin}} - \overrightarrow{\text{Germany}} + \overrightarrow{\text{Japan}} = \overrightarrow{\text{Tokyo}}$ (Mikolov et al. 2013)
- GloVe: 80% accuracy!

Linear relations between countries and capitals in GloVe



Google Analogy Test (Mikolov et al 2013)

9 morphological categories: adjective-to-adverb, comparatives, superlatives, verb:present-participle, country-nationality, verb:past-tense, verb:3PsSg-plural, opposites.

5 semantic categories: common countries and capitals, countries and capitals of the world, city-in-state, country-and-currency, male:female.

20-70 unique word pairs per category.

8,869 semantic and 10,675 morphological questions in total.

Bigger Analogy Test (Gladkova et al 2016)

Inflectional morphology	Nouns	regular plurals (<i>student:students</i>), plurals with orthographic changes (<i>wife:wives</i>)
	Adjectives	comparative degree (<i>strong:stronger</i>), superlative degree (<i>strong:strongest</i>)
	Verbs	infinitive: 3Ps.Sg (<i>follow:follows</i>), infinitive: participle (<i>follow:following</i>), infinitive: past (<i>follow:followed</i>), participle: 3Ps.Sg (<i>following:follows</i>), participle: past (<i>following:followed</i>), 3Ps.Sg : past (<i>follows:followed</i>)
Derivational morphology	Stem change	verb+er (<i>bake:baker</i>), verb+able (<i>edit:editable</i>), verb+ation (<i>continue:continuation</i>), verb+ment (<i>argue:argument</i>)
	No stem change	re+verb (<i>create:recreate</i>), noun+less (<i>home:homeless</i>), adj.+ness (<i>mad:madness</i>), un+adj. (<i>able:unable</i>), adj.+ly (<i>usual:usually</i>), over+adj. (<i>used:overused</i>)
Lexicographic semantics	Hypernyms	animals (<i>turtle:reptile</i>), miscellaneous (<i>peach:fruit</i>)
	Hyponyms	miscellaneous (<i>color:white</i>)
	Meronyms	part-whole (<i>car:engine</i>), substance (<i>sea:water</i>), member (<i>player:team</i>),
	Antonyms	opposites (<i>up:down</i>), gradable (<i>clean:dirty</i>)
	Synonyms	exact (<i>sofa:couch</i>), intensity (<i>cry:scream</i>)
Encyclopedic semantics	Animals	the young (<i>cat:kitten</i>), sounds (<i>dog:bark</i>), shelter <i>fox:den</i>
	Geography	capitals (<i>Athens:Greece</i>), languages (<i>Peru:Spanish</i>), UK city:county <i>York:Yorkshire</i>
	People	occupation (<i>Lincoln:president</i>), nationalities (<i>Lincoln:American</i>)
	Other	thing:color (<i>blood:red</i>), male:female (<i>actor:actress</i>)

Analogies?

Why should different linguistic relations translate to exactly the same vector offsets for all words?

Other Methods of Evaluation

- Extrinsic evaluation via performance on downstream tasks** (POS-tagging, chunking, NER, sentiment polarity, NLI)
- Behavioral evaluation:** correlation with similarity judgments, intrusion, N400 effect, fMRI scans, eye-tracking, and semantic priming data.

Benefits of Neural Approaches

Easy to learn, especially with good linear algebra libraries.

Highly parallel problem: minibatching, GPUs, distributed models.

Can predict other discrete aspects of context (dependencies, POS tags, etc). Can estimate these probabilities with counts, but sparsity quickly becomes a problem.

Can predict/condition on continuous contexts: e.g. images.

Comparison with count-based methods

Count based and objective-based models: same general idea.

Word2Vec == PMI matrix factorization of count based models (Levy and Goldberg, 2014)

Count-based and most neural models have equivalent performance when properly hyper parameters are properly optimized (Levy et al. 2015)



Tokenization

Remember tokenization?

Tokenization is splitting text into meaningful units that form your vocabulary

- words, punctuation
- can use white-space segmentation to get words (roughly)

Tokens are not just words

Word-internal punctuation: Ph.D., AT&T, Google.com, 555,500.50

Expanding clitics: I'm → I am

Multiword tokens: New York, Rock 'n' roll

Word vectors are really token vectors!

Out-of-vocabulary words

Texts are sparse!

Many words (tokens) never seen even in large texts.

What if your test data contains words that are not in your training data?

You get the so-called “out of vocabulary” words (OOV)

People would add a special “<unknown>” token to their vocabulary

During training, use this token for unseen words in validation data.

A better solution – subword tokenization

- Can we construct a vocabulary of meaningful subwords?
- Use the data to tell us how to tokenize!
- Tokens could be subwords, i.e. variable-length character ngrams
- That way, out-of-vocabulary words can sometimes be represented reasonably.

Subword tokenization

- Byte-Pair Encoding (Sennrich et al, 2016)
- WordPiece (Schuster and Nakajima, 2012)

Two parts:

Token learner takes a raw training corpus and induces a vocabulary (a set of tokens)

Token segmenter takes a raw test corpus and tokenizes it according to the vocabulary



Byte Pair Encoding (BPE) Tokenization

- Let the initial vocabulary be the set of individual characters = {A, B, C, D, ... , a, b, c,}

Repeat

- Choose two symbols that are most often adjacent in the training corpus (e.g. “t” and “h”)
- Add a new merged symbol “th” to the vocabulary
- Replace every adjacent “t” and “h” with “th” in the training corpus

Until k merged have been done.

Byte Pair Encoding (BPE) Tokenization

- BPE will often deduce frequent subwords: morphemes like –est, –er, un–, etc.
- Most subword algorithms are run inside space-separated tokens, adding a special end-of-word character before each space.
- That way, word-final combinations of letters get a different treatment than word-internal combinations

BPE Example

Training text: “This runner jumped higher”

- Initial vocabulary:

t, h, i, s, r, u, n, e, j, u, m, d, j, i, g, <eow>

- Initial tokenization

this <eow> runner <eow> jumped <eow> higher <eow>

- Merge “e” and “r”:

This <eow> run **er** <eow> jumped <eow> high **er** <eow>

- Merge “er” and “<eow>

This <eow> runn **er<eow>** jumped <eow> high **er<eow>**

WordPiece

- Popularized by BERT, which was the first pretrained transformer encoder.
- Instead of relying on the frequency of token pairs during the merge, at each merge step:
 - Train a language model at each step
 - Merge the token that maximizes the likelihood of the training data, i.e.
 - $p(t_1 t_2) > p(t_1) p(t_2)$
 - $MI(t_1, t_2)$ is greater than for any other pair



Text Classification

Text Classification Examples

- Identifying topics
 - sports / politics / finance / ...
- Subject headings for books / articles
 - MeSH headings: Drug Therapy / Embryology / ...
- Sentiment analysis
- Authorship identification
- Spam detection
- etc.



Text Classification

Input:

a document $d \in D$

a fixed set of classes $C = \{c_1 \dots c_k\}$

Output:

a predicted class $c \in C$

a trained model \rightarrow a learned classification function $f : D \rightarrow C$

Workflow

- Create a **labeled corpus**
select texts, pick categories, assign labels
- Split it into **test** and **training** segments

Supervised
Task

- Choose a representation, i.e. how each text will be represented
 - Choose a classifier model
- Naive Bayes, Decision Tree, Logistic Regression, SVM, Neural Network



- Train the model on the training data
- Test model performance on the test data
- If satisfactory, use the model to process unseen text

Text Classification Models

Could be any of your favorite classifier models:

Logistic regression (LR), support vector machine (SVM), multi-layer perceptron (MLP), deep neural network (the last hidden state is used as input to a softmax layer that does classification), etc.

Consider text classification done by applying **logistic regression** (softmax) to **vector representations of input text**

Project into an softmax layer of dimensionality $|C|$

Apply softmax to find highest-probability category

Multiclass Logistic Regression

Vector representation of the document is used as input to a softmax layer that does classification, assigning a probability to each label y_i :

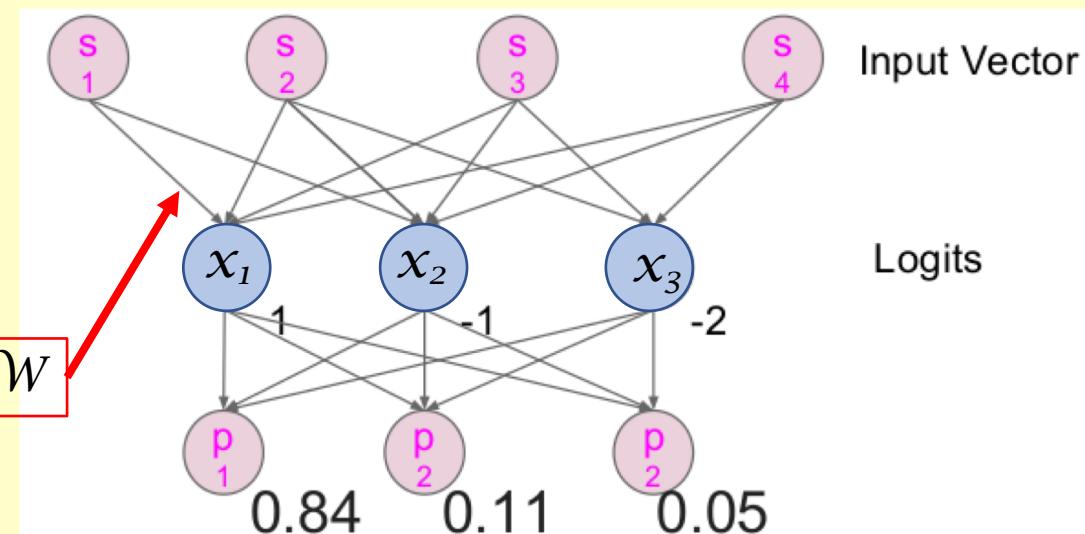
$$P(y = j|x) = \frac{e^{x_j}}{\sum_{k=1}^K e^{x_k}}$$

Weight matrix \mathcal{W}

$$\mathbf{x} = \mathcal{W} \mathbf{s}$$

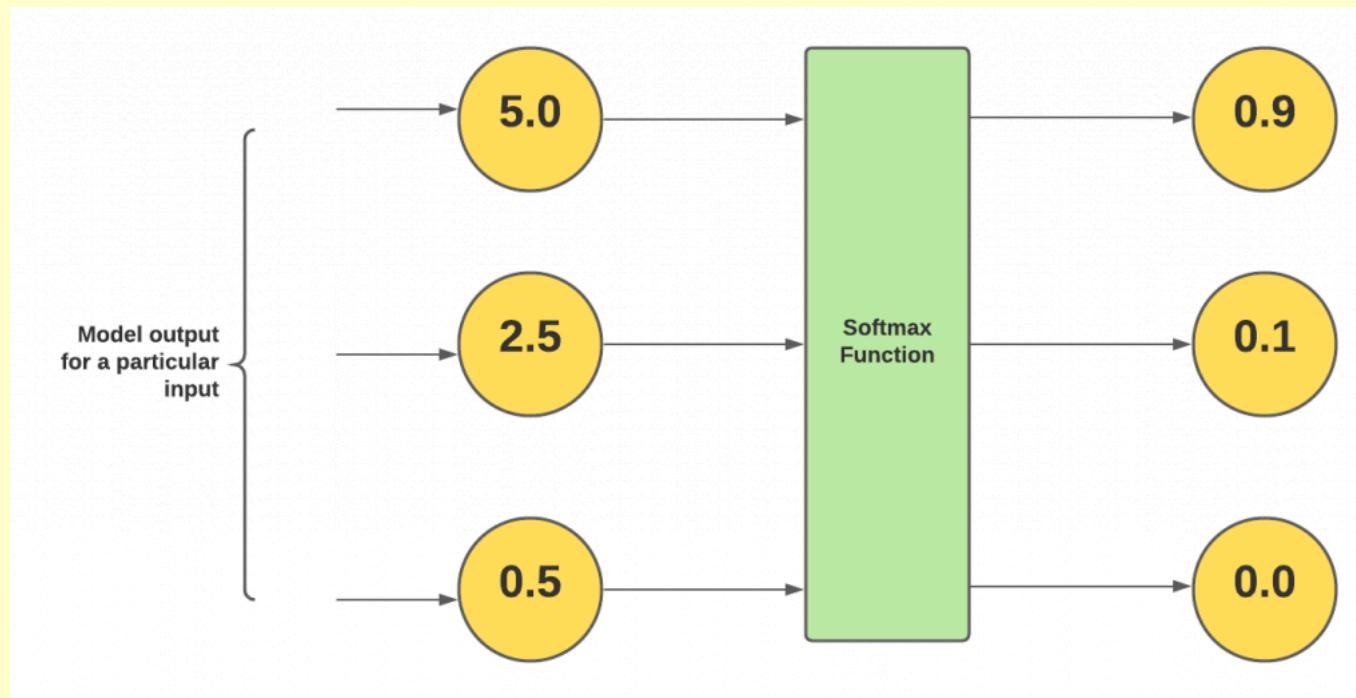
$$x_i = \mathbf{w}_j \mathbf{s}$$

$$x_i = \sum s_j w_{ij}$$



Softmax operation

$$P(y = j|x) = \frac{e^{x_j}}{\sum_{k=1}^K e^{x_k}}$$



Vector Space Embeddings

“Count-based” vectors

Variously normalized word count-based representations

Vocabulary-size dimensionality reduced via PCA, SVD, etc.

Learned vectors (“prediction-based”)

Learn embeddings as parameters in a learning task

Where the cost function e.g. maximizes the probability of your training text

Bag-of-Words (BOW) Representation

Element-wise “add” of count-based word vectors

Counts could be normalized by frequency

Probability (divide by corpus size)

Conditional probability (divide by co-occurrence count)

Pointwise mutual information

TF-IDF

Dimensionality reduction via PCA, SVD, etc.

Learned vectors

Prediction-based (i.e. learned) vector space embeddings

Word-level embeddings

Text-level embeddings (sentence, paragraph, etc.)

Parameter vectors, trained to maximize e.g.,

e.g. the probability of word sequences in available text

Learned vectors

Using learned **word-level embeddings** for sentence representation

How do we use them?

What's the “BoW” representation?



Baselines for Representing Input Text

BoW approach with *learned* word vectors

Element-wise “add” of learned word vectors, equivalent to:

Mean-pooling (average) over pretrained word vectors (learned in pre-training on free text)

May also do max-pooling (max)

Baselines for Representing Input Text

“A simple but tough to beat baseline for sentence embeddings”. Arora et al ICLR 2017.

Compute sentence representations using weighted averaging over BoW word embeddings

$a/(a + p(w))$ with a being a parameter and $p(w)$ the (estimated) word frequency

- Similar to TF-IDF, downweigh frequent words

Do PCA (principal component analysis) over sentences

Subtract from the sentence representation its projection on the first PCA component

Good unsupervised performance on semantic similarity/relatedness; good initialization on supervised classification tasks (sentiment, entailment, etc.)

Sample Classification Tasks

Sentiment

“This movie was actually neither that funny, nor super witty.”

Labels: **positive**, **negative**, neutral

Entailment

“The maniac killed his victim” – “His victim died”

“The maniac killed his victim” – “His victim survived”

Labels: **Entailment**, **Contradiction**, Neutral

Learned vectors

Prediction-based (i.e. learned) vector space embeddings

Word-level embeddings

Text-level embeddings (sentence, paragraph, etc.)

Parameter vectors, trained to maximize e.g.,

e.g. the probability of word sequences in available text

Learned vectors

Prediction-based (i.e. learned) vector space embeddings

Word-level embeddings

Text-level embeddings (sentence, paragraph, etc.)

Parameter vectors, trained to maximize e.g.,

e.g. the probability of word sequences in available text

Semi-supervised objective

Vector space embeddings for linguistic inputs in DL are usually *learned* by using some semi-supervised objective function

i.e. the function that does not require manually tagged data

For example, a function that **maximizes the probability of existing data**

As we saw in learned word-level embeddings (lexical vectors).

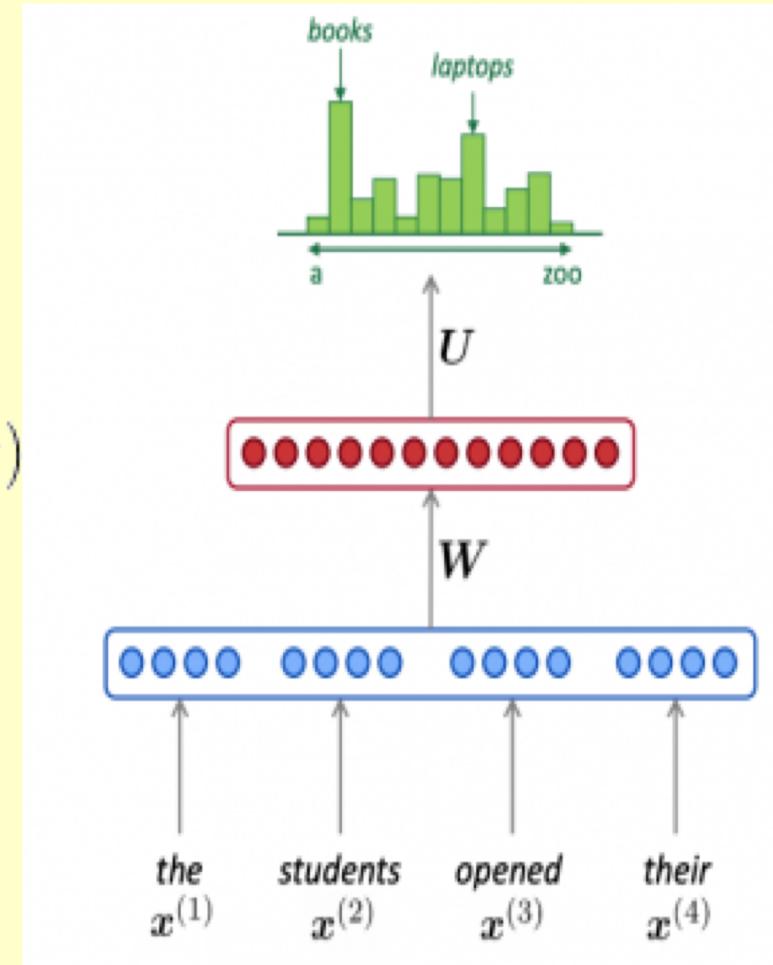
Language Modeling

Maximizing the probability of existing data

This is the so-called language modeling:
assigning probabilities to word sequences
(or character sequences).

$$\begin{aligned} P(w_1^t) &= P(w_1)P(w_2|w_1)P(w_3|w_1w_2)\dots P(w_t|w_1^{t-1}) \\ &= \prod_{i=1}^t P(w_i|w_1^{i-1}) \end{aligned}$$

There is unlimited data on which to train these models: *the model just needs to learn to predict the words in existing text.*



Semi-supervised objectives

There are other semi-supervised objectives, as we'll see.

- Masked language model (denoising auto-encoder objective)
- Span in-filling objective

Sentence-based objectives

- predicting next sentence
- ordering sentences
- discriminating between next sentence and random sentence, etc.

Prediction-based Vectors

ConvNets (CNNs) applied over concatenated word vectors to do classification

[popular ca. 2013 – 2015]

Skip-thought sentence representations based on reconstructing previous / next sentence [2015]

GRU RNN encoder/decoder

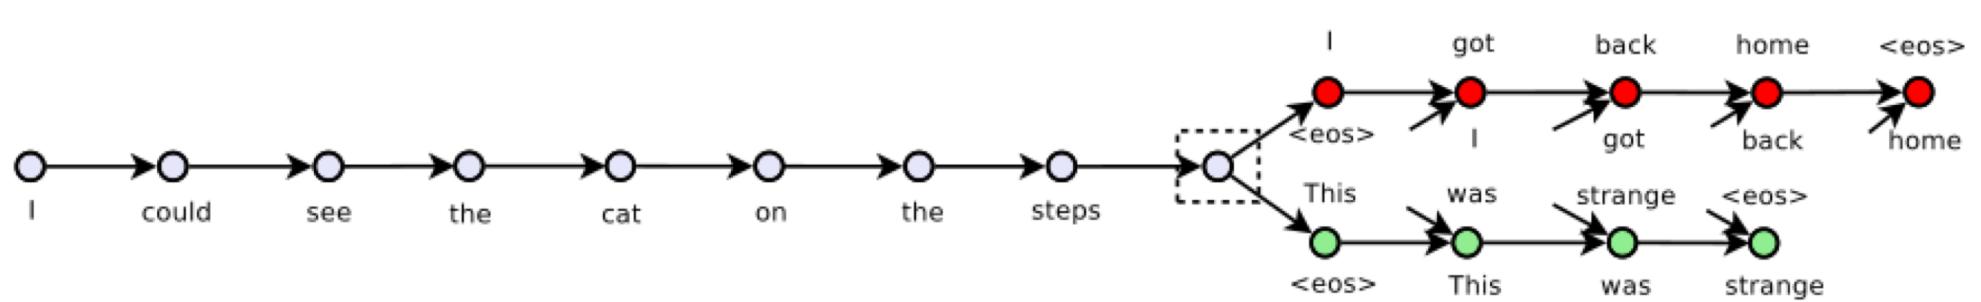


Figure 1: The skip-thoughts model. Given a tuple (s_{i-1}, s_i, s_{i+1}) of contiguous sentences, with s_i the i -th sentence of a book, the sentence s_i is encoded and tries to reconstruct the previous sentence s_{i-1} and next sentence s_{i+1} . In this example, the input is the sentence triplet *I got back home. I could see the cat on the steps. This was strange.* Unattached arrows are connected to the encoder output. Colors indicate which components share parameters. $\langle \text{eos} \rangle$ is the end of sentence token.

Prediction-based Vectors

Sequence-to-sequence autoencoders – optimizes the reconstruction accuracy for the input sentence, which is reconstructed word by word.

Special **CLS embeddings** learned by attention-based Transformer models trained to predict masked words in free text

BERT CLS embeddings (pre-trained bi-directional transformer)
[Devlin et al 2018]

Universal Sentence Encoder (Transformer trained on a skip-thought sentence-prediction task, and several supervised tasks) [Cer et al 2018]

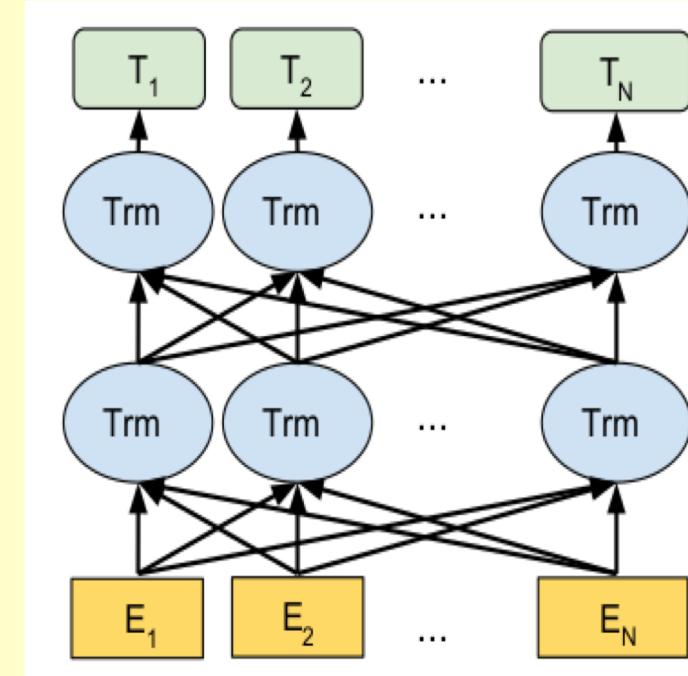
Transformer-Based Pre-training for Input Representations

Special CLS embeddings learned by attention-based Transformer models are used to represent input text

BERT [Devlin et al 2019], a Transformer-based encoder model is trained to predict a) masked words in free text (“masked language modeling” objective, and b) next sentence

BERT (and similar architectures) are trained on free text (“**pre-training**”), then **fine-tuned** for the specific task (e.g. text classification), using supervised learning.

Input tokens are typically subwords, and use [WordPiece](#) [Schuster and Nakajima 2012] or (more commonly) [Byte Pair Encoding](#) [Sennrich et al. 2016] to deal with out-of-vocabulary problem.



BERT Encoder



Lab : Homework 1 (see course schedule)